



WeCloudData

# Instacart Market Basket Analysis


Allen XU

# Agenda

- Introduction
- Business Goal
- Exploratory Data Analysis
- Feature Engineering
- Model Select
- Model Interpretation
- Apply Model
- Challenge & Next Step


# Intro ML For Retail

- Personalized Marketing ★
- Demand Forecasting
- Supply Chain
- Customer Experience



# Personalized Marketing

- Others You May Like
- Frequently Bought Together (shopping cart expansion)
- Recommended for You
- Similar Items
- **Buy it Again** ★
- On-sale / Deals
- Recently Viewed



# Goal

- Set Up Business Goal
  - Simplify user purchase journey (short term)
  - improve user retention rate (long term)
- Define Machine Learning Problem
  - Binary classification: Will Reorder/Won't Reorder
  - Statistic Metric
    - F1
    - Accuracy



# Data

```
order_products_train_df.head()
```

	order_id	product_id	add_to_cart_order	reordered
0	1	49302	1	1
1	1	11109	2	1
2	1	10246	3	0
3	1	49683	4	0
4	1	43633	5	1

```
order_products_prior_df.head()
```

	order_id	product_id	add_to_cart_order	reordered
0	2	33120	1	1
1	2	28985	2	1
2	2	9327	3	0
3	2	45918	4	1
4	2	30035	5	0

```
aisles_df.head()
```

	aisle_id	aisle
0	1	prepared soups salads
1	2	specialty cheeses
2	3	energy granola bars
3	4	instant foods
4	5	marinades meat preparation

```
aisles_df.shape
```

```
(134, 2)
```

```
departments_df.head()
```

	department_id	department
0	1	frozen
1	2	other
2	3	bakery
3	4	produce
4	5	alcohol

```
orders_df.head()
```


	order_id	user_id	eval_set	order_number	order_dow	order_hour_of_day	days_since_prior_order
0	2539329	1	prior	1	2	8	NaN
1	2398795	1	prior	2	3	7	15.0
2	473747	1	prior	3	3	12	21.0
3	2254736	1	prior	4	4	7	29.0
4	431534	1	prior	5	4	15	28.0

```
products_df.head()
```


	product_id	product_name	aisle_id	department_id
0	1	Chocolate Sandwich Cookies	61	19
1	2	All-Seasons Salt	104	13
2	3	Robust Golden Unsweetened Oolong Tea	94	7
3	4	Smart Ones Classic Favorites Mini Rigatoni Wit...	38	1
4	5	Green Chile Anytime Sauce	5	13

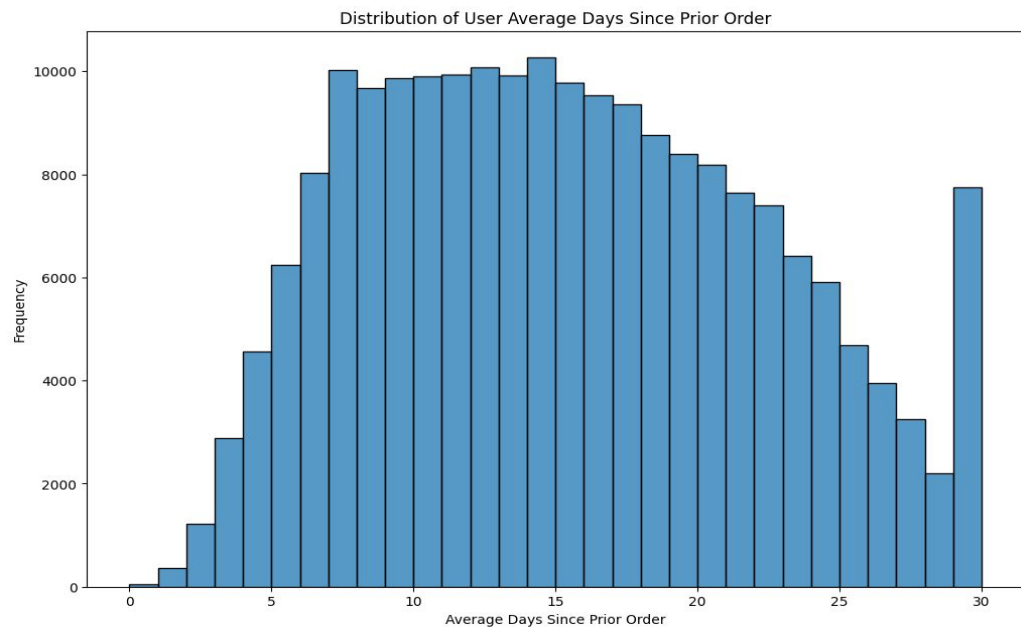
# Data Exploration

Top 10 Products with Highest Reorder Rate



Top 20 Aisles by Reorder Rate







# Feature Engineering

## User Part

```
# _user_total_orders:
# _user_sum_days_since_prior_order:
# _user_mean_days_since_prior_order:
agg_dict_2 = {'order_number': [('_user_total_orders', 'max')],
              'days_since_prior_order': [('_user_sum_days_since_prior_order', 'sum'),
                                           ('_user_mean_days_since_prior_order', 'mean')]}

users = ka_add_groupby_features_1_vs_n(orders[orders.eval_set == 'prior'], ['user_id'], agg_dict_2)

# _user_reorder_ratio: total number of reorder/total number of orders of first order
# _user_total_products:
# _user_distinct_products:
agg_dict_3 = {'reordered': [('_user_reorder_ratio',
                             lambda x: sum(priors_orders_detail.loc[x.index, 'reordered'] == 1) /
                             sum(priors_orders_detail.loc[x.index, 'order_number'] > 1))],
              'product_id': [('_user_total_products', 'count'),
                             ('_user_distinct_products', lambda x: x.nunique()))]

us = ka_add_groupby_features_1_vs_n(priors_orders_detail, ['user_id'], agg_dict_3)
users = users.merge(us, how='inner')

# average number of item in each orders
users['_user_average_basket'] = users['_user_total_products'] / users['_user_total_orders']

us = orders[orders.eval_set != "prior"][['user_id', 'order_id', 'eval_set', 'days_since_prior_order']]
us.rename(columns={'days_since_prior_order': 'time_since_last_order'}, inplace=True)

users = users.merge(us, how='inner')
```

## Item part

```
# add order information to priors set
priors_orders_detail = orders.merge(right=priors, how='inner', on='order_id')

# create new variables
# _user_buy_product_times: the time user by this item
priors_orders_detail['_user_buy_product_times'] = priors_orders_detail.groupby(['user_id', 'product_id']).cumcount() + 1

# _prod_tot_cnts: total number that item been sold
# _reorder_tot_cnts_of_this_prod: reorder number of this item
# _prod_order_once: the number of times that item only buy 1
# _prod_order_more_than_once: the number of times that item buy more than once
agg_dict = {'order_id': [('_total_orders', 'count')],
            'add_to_cart_order': [('_sum_add_to_cart_order', 'sum')],
            'reordered': [('_reorder_total_cnt', 'sum')],
            '_user_buy_product_times': [('_prod_buy_first_time_total_cnt', lambda x: sum(x == 1)),
                                         ('_prod_buy_second_time_total_cnt', lambda x: sum(x == 2))]}

prd = ka_add_groupby_features_1_vs_n(priors_orders_detail, ['product_id'], agg_dict)

# _prod_reorder_prob:
# _prod_reorder_ratio:
prd['_prod_reorder_prob'] = prd['_prod_buy_second_time_total_cnt'] / prd['_prod_buy_first_time_total_cnt']
prd['_prod_reorder_ratio'] = prd['_reorder_total_cnt'] / prd['_total_orders']
prd['_prod_reorder_times'] = 1 + prd['_reorder_total_cnt'] / prd['_prod_buy_first_time_total_cnt']
```

## General Part

```
# _up_order_count: number of times user buy this item
# _up_first_order_number: order number of user that first time buy this item
# _up_last_order_number: order number of user that last time buy this item
# _up_average_cart_position:
agg_dict_4 = {'order_number': [('_up_order_count', 'count'),
                               ('_up_first_order_number', 'min'),
                               ('_up_last_order_number', 'max')],
              'add_to_cart_order': [('_up_average_cart_position', 'mean')]}

data = ka_add_groupby_features_1_vs_n(df=priors_orders_detail,
                                     group_columns_list=['user_id', 'product_id'],
                                     agg_dict=agg_dict_4)

data = data.merge(prd, how='inner', on='product_id').merge(users, how='inner', on='user_id')
# _up_order_rate : number of times buy this item/ total order number
# _up_order_since_last_order
# _up_order_rate_since_first_order : number of purchase of this item/ the number of orders between first time buy this item and last time buy this item
data['_up_order_rate'] = data['_up_order_count'] / data['_user_total_orders']
data['_up_order_since_last_order'] = data['_user_total_orders'] - data['_up_last_order_number']
data['_up_order_rate_since_first_order'] = data['_up_order_count'] / (data['_user_total_orders'] - data['_up_first_order_number'] + 1)

# add user_id to train set
train = train.merge(right=orders[['order_id', 'user_id']], how='left', on='order_id')
data = data.merge(train[['user_id', 'product_id', 'reordered']], on=['user_id', 'product_id'], how='left')

# release memory
del priors_orders_detail, orders
gc.collect()
data.head()
```

# Model Select (Metrics F1&Accuracy)

## LR

*# Initialize a Logistic Regression model with specific*

```
lr_model = LogisticRegression(n_jobs=16,  
                              random_state=42,  
                              class_weight='balanced')
```

*# Fit the model using the training data*

```
lr_model.fit(X_train, y_train)
```

F1 Score: 0.2246

Accuracy Score: 0.7964

## RF

*# Create a Random Forest model instance*

```
rf_model = RandomForestClassifier(n_jobs=16,  
                                 random_state=42)
```

*# Fit the model using the training data*

```
rf_model.fit(X_train, y_train)
```

F1 Score: 0.2693

Accuracy Score: 0.9085

## XGB

```
xgb_params = {  
    "objective"      : "reg:logistic"  
    , "eval_metric"  : "logloss"  
    , "eta"          : 0.1  
    , "max_depth"    : 6  
    , "min_child_weight" : 10  
    , "gamma"        : 0.70  
    , "subsample"     : 0.76  
    , "colsample_bytree" : 0.95  
    , "alpha"        : 2e-05  
    , "lambda"       : 10  
    , "n_jobs"       : 16  
}
```

```
watchlist = [(d_train, "train")]  
bst = xgboost.train(params=xgb_params,  
                    dtrain=d_train,  
                    num_boost_round=80,  
                    evals=watchlist, verbose_eval=10)
```

F1 Score: 0.4419

Accuracy Score: 0.8747

## DT

*# Train a Decision Tree classifier with specified parameters*

```
dt = DecisionTreeClassifier(max_depth=10,  
                            min_samples_split=10,  
                            min_samples_leaf=10,  
                            class_weight='balanced')
```

```
dt.fit(X_train, y_train)
```

F1 Score: 0.3647

Accuracy Score: 0.7481

# Model Select(Kaggle Results)

## XGB

order_id		products	
0	17	13107 21463	
1	34	21137 39180 47792 43504 13176 47766 44663 1608...	
2	137	38689 5134 41787 23794 2326 25890 24852	
3	182	32109 41149 33000 47672 39275 9337 11520 47209...	
4	257	21137 27966 28476 49235 29837 37646 27104 4605...	
...	...	...	
74995	3420740	19660 13176 16291 49005 8174 35951 39146 46252	
74996	3420877	21137 27966 43122 10768 13176 49111 12238 8277...	
74997	3420888	20542 43961 44632 22935 35951 46906 28985 4776...	
74998	3420989	5134 46676 47229 21616 42450 35004 43210 47766...	
74999	3421054	18426 38448 11123 31231 13375 24852	

75000 rows × 2 columns

product_id	product_name	
13106	13107	The Complete Cookie Double Chocolate
21462	21463	Vanilla Bean Light Ice Cream

## DT

order_id		products	
0	17	31964 13107 21903 39275 39928 21463 38777...	
1	34	21137 7948 13176 30353 49191 2596 39180 47792 ...	
2	137	38689 24235 43352 29594 23794 25890 34126 5134...	
3	182	32109 28800 22935 15252 47672 13629 40377 3927...	
4	257	21137 38693 24838 30233 24964 27966 29837 2203...	
...	...	...	
74995	3420740	19660 11777 38772 46252 5785 16291 39928 10017...	
74996	3420877	21137 11005 5449 6567 5918 7781 13176 12238 26...	
74997	3420888	31506 22935 28985 42701 24964 30450 42123 4396...	
74998	3420989	45866 7781 42450 28985 30406 27521 5134 21616 ...	
74999	3421054	27429 18426 5818 38448 27370 11123 15732 15802...	

75000 rows × 2 columns

product_id	product_name	
13106	13107	The Complete Cookie Double Chocolate
21462	21463	Vanilla Bean Light Ice Cream
21902	21903	Organic Baby Spinach
31963	31964	Passionfruit Sparkling Water
38776	38777	Organic Green Seedless Grapes
39274	39275	Organic Blueberries
39927	39928	Organic Kiwi

## LR

order_id		products	
0	17	16965 21709 15613 21903 39275 39928 47766 3877...	
1	34	21137 13176 39275 47766 46979 42265 44632	
2	137	26209 24184 18465 41787 21903 43352 16797 2710...	
3	182	31717 17794 21903 13176 27104 44359 39275 2293...	
4	257	21137 13176 24838 28985 30233 42701 24964 2796...	
...	...	...	
74995	3420740	21137 19660 13176 43352 26604 22935 5077 24964...	
74996	3420877	21137 13176 26604 22935 24964 47626 27966 4605...	
74997	3420888	8193 8424 28204 42736 43961 44632 31506 22935 ...	
74998	3420989	27521 7781 46676 17948 31506 21616 8277 35951 ...	
74999	3421054	28842 24852	

75000 rows × 2 columns

product_id	product_name	
15612	15613	Salted Tub of Butter
16964	16965	Chocolate Ice Cream
21708	21709	Sparkling Lemon Water
21902	21903	Organic Baby Spinach
38776	38777	Organic Green Seedless Grapes
39274	39275	Organic Blueberries
39927	39928	Organic Kiwi
47765	47766	Organic Avocado

## RF

order_id		products
0	17	13107
1	34	16083
2	137	None
3	182	9337 13629 39275
4	257	49235
...	...	...
74995	3420740	None
74996	3420877	21137 13176 13646
74997	3420888	35951 44632
74998	3420989	47766
74999	3421054	11123 31231



75000 rows × 2 columns

product_id	product_name	
13106	13107	The Complete Cookie Double Chocolate

# Deep Learning Model

- MLPClassifier
- 'sgd' refers to stochastic gradient descent.
- Adam solver with early\_stopping to prevent overfit

```
model = MLPClassifier(solver='adam',  
                      activation='relu',  
                      learning_rate_init = 0.0000001,  
                      alpha=1e-5,  
                      hidden_layer_sizes=(128, 32, 16),  
                      random_state=1,  
                      warm_start=True,  
                      batch_size = 4096,  
                      verbose = True,  
                      tol=1e-10,  
                      early_stopping=True)  
  
model.fit(X_train.values, y_train.values)
```



# Ensemble Methods

```
xgb_params = {
    "objective" : "reg:logistic"
    , "eval_metric" : "logloss"
    , "eta" : 0.1
    , "max_depth" : 6
    , "min_child_weight" : 10
    , "gamma" : 0.70
    , "subsample" : 0.76
    , "colsample_bytree" : 0.95
    , "alpha" : 1e-05
    , "lambda" : 10
    , "n_jobs" : 16
}

# Create an XGBoost model
xgb_model = xgboost.XGBClassifier(**xgb_params)

# Create a Decision Tree model
dt_model = DecisionTreeClassifier(max_depth=10, min_samples_split=10, min_samples_leaf=10)

# Create a Logistic Regression model
lr_model = LogisticRegression(n_jobs=16, random_state=42)

# Create a Voting Classifier instance, including all models
voting_classifier = VotingClassifier(estimators=[
    ('xgb', xgb_model),
    ('dt', dt_model),
    ('lr', lr_model)
], voting='hard')

# Fit the model using the training data
voting_classifier.fit(X_train, y_train)

# Make predictions on the test data
X_test.loc[:, 'reordered'] = voting_classifier.predict(X_test.drop(['eval_set', 'user_id', 'order_id', 'reordered', 'product_id'], axis=1))

# Keep the same post-processing steps as before
X_test.loc[:, 'product_id'] = X_test.product_id.astype(str)

submit = ka_add_groupby_features_n_vs_1(X_test[X_test.reordered == 1],
                                       group_columns_list=['order_id'],
                                       target_columns_list=['product_id'],
                                       methods_list=[lambda x: ' '.join(set(x))], keep_only_stats=True)

submit.columns = sample_submission.columns.tolist()

submit_final_hard = sample_submission[['order_id']].merge(submit, how='left').fillna('None')
submit_final_hard
```

Hard

order_id		products					
0	17	38777 21903 39928 39275					
1	34	21137 13176 47766 42265 44632					
2	137	26209 24184 18465 41787 21903 43352 27104 3412...					
3	182	39275 22935 9337 47209 5479 8518					
4	257	21137 12341 27966 49235 37646 22035 27104 4605...					
...	...	...					
74995	3420740	19660 48679 8174 39928 35951					
74996	3420877	21137 27966 21903 13176 26604 8277 47209 40706...					
74997	3420888	8193 8424 42736 43961 44632 31506 22935 35951 ...					
74998	3420989	27521 7781 46676 17948 21616 8277 35951 47766 ...					
74999	3421054	11123 31231 24852					

75000 rows × 2 columns

Soft


order_id		products							
0	17	21903 13107 21463 39275							
1	34	21137 13176 47766 16083 2596 39475							
2	137	38689 5134 41787 21903 23794 2326 25890 24852							
3	182	33000 47672 39275 9337 11520 47209 5479 13629							
4	257	21137 27966 49235 29837 37646 27104 4605 45013...							
...	...	...							
74995	3420740	19660 49005 39928 35951 39146							
74996	3420877	21137 27966 43122 21903 13176 49111 12238 8277...							
74997	3420888	43961 44632 22935 7963 35951 46906 47766 28985...							
74998	3420989	46676 21616 42450 35004 43210 47766 13517							
74999	3421054	18426 11123 31231 13375 24852							

75000 rows × 2 columns




# Feature Importance


XGB Classifier Feature Importance




Decision Tree Classifier Feature Importance



Feature Importance in Logistic Regression Model



Feature Importance in Random Forest Model



# Apply Model

- Perform A/B test before launch
- After launch, continue to monitor model performance (live data support) & business performance
- After launch, continue to collect live traffic data to improve model

# Challenges & Improvement

## Challenges

- Big Data Low Ram
- Competition Problem
- Wrong Direction

## Improvement

- Thinking More
- Add Better Features
- Try More DL Model
- Try More Ensemble Method
- Keep Study on ML